

NaluNet

Darknet + YOLO v3 + Python/Flask

By @pjcr and @lemiffe

Special thanks to @jessedobbelaere, @bugragokalp, @tmeire_, @pjreddie, @wonko_be

—— THE PROBLEM ——

IS THERE ANY NALU?

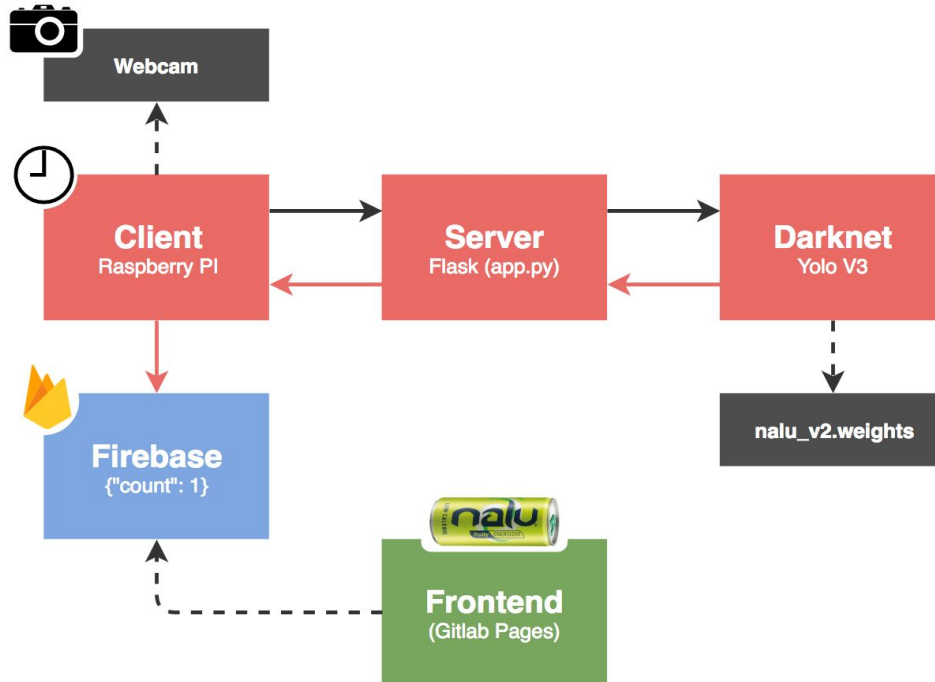
THE GOAL

CREATE A **NEURAL NETWORK** THAT CAN DETECT THE **PRESENCE OF CANS OF NALU** AND UPDATE A WEBSITE WITH THE ESTIMATED COUNT EVERY ~10 SECONDS.

(YES I KNOW, WE ARE LAZY :D WE COULD JUST STAND UP, WALK FOR A BIT AND CHECK, BUT THIS IS WAY MORE FUN)



THIS IS HOW IT WORKS...



— STEP 1 —

TRAINING THE AI

FIRST WE TOOK OVER 100 PHOTOS



home / Nalu

Nalu Detection System

[START LABELING](#)

You've labeled everything available to you. Other labelers must finish the remaining labeling tasks.

[OVERVIEW](#)[ACTIVITY](#)[PERFORMANCE](#)[EXPORT](#)[SETTINGS](#)

Progress

155

Submitted

70

Remaining


17

Skipped

71%

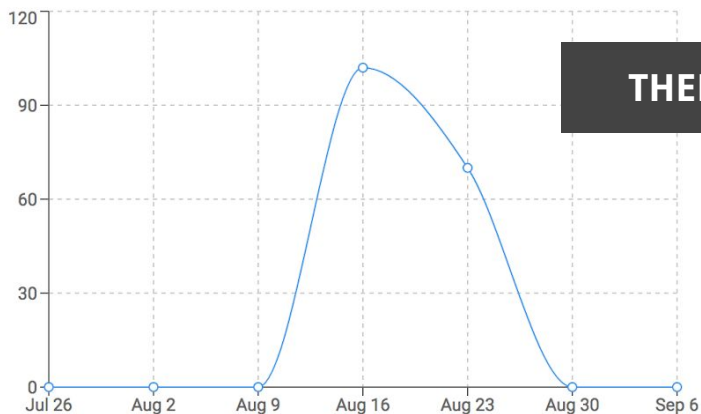
Complete

Object Count ▾



Labeling Activity

Weekly ▾



THEN WE PUT BOUNDARIES AROUND EACH NALU CAN


```

1 import pandas as pd
2 from shutil import copyfile
3 from PIL import Image, ImageFont, ImageDraw, ImageEnhance
4 from os import path
5
6 boundingBoxesFile = 'nalu.json'
7 picturesFolder = 'Nalu_Pictures'
8 prefix = 'data'
9 outputFolder = 'obj'
10 checkFolder = 'check'
11
12 df = pd.read_json(boundingBoxesFile, orient='values')
13
14 def check():
15     pass
16
17 def parseBounds(coordsArr, img_height, img_width, label):
18     x_l = coordsArr[0]['x']
19     y_l = coordsArr[0]['y']
20     x_h = coordsArr[2]['x']
21     y_h = coordsArr[2]['y']
22
23     x_avg = (x_h + x_l)/2.
24     y_avg = (y_h + y_l)/2.
25     width = abs(x_h - x_l)
26     height = abs(y_h - y_l)
27
28     return [label, x_avg/img_width, 1. - y_avg/img_height, width/img_width, height/img_height]
29
30 with open('train.txt', 'w') as train_file:
31     for i, row in enumerate(df.iterrows()):
32         tmpDict = row[1]['Label']
33         if(type(tmpDict) == dict):
34             image_name = row[1]['External ID']
35             train_file.write(path.join(prefix, outputFolder, image_name) + '\n')
36             image = path.join(picturesFolder, image_name)
37             im = Image.open(image)
38             width, height = im.size
39
40             copyfile(path.join(picturesFolder, image_name), path.join(outputFolder, image_name))
41             source_img = Image.open(image)
42             draw = ImageDraw.Draw(source_img)
43             with open(path.join(outputFolder, image_name.split('.')[0] + '.txt'), 'w') as f:
44                 for box in tmpDict['Nalu']:
45                     labelProps = parseBounds(box['geometry'], height, width, 0)
46                     x1 = (labelProps[1] - labelProps[3] / 2.) * width
47                     y1 = (labelProps[2] - labelProps[4] / 2.) * height
48                     x2 = (labelProps[1] + labelProps[3] / 2.) * width
49                     y2 = (labelProps[2] + labelProps[4] / 2.) * height
50                     draw.rectangle((x1, y1), (x2, y2)), fill="black")
51                     f.write(' '.join(str(x) for x in labelProps) + '\n')
52             source_img.save(path.join(checkFolder, image_name), "JPEG")
53

```

DATA PREPARATION

THEN WE **EXPORTED THE FILE AS**
A .JSON, PLUS ALL THE PHOTOS.

AFTERWARDS WE PERFORMED
SEVERAL **TRANSFORMATIONS**
ON THE RESULTS AND EXPORTED
THE DATA TO A **TXT FORMAT**
THAT DARKNET CAN INTERPRET.

WE HAD TO **INVERT THE COORDS**
AS DARKNET/YOLO READ THE Y
AXIS AS STARTING FROM BELOW.

AI TRAINING SERVER



nalu

us-east1-c

(nic0)

SSH



WE THEN TRAINED THE MODEL USING DARKNET/YOLO ON A SERVER:

8 vCPUs, 30 GB RAM

1 x NVIDIA Tesla K80

WHAT IS DARKNET & YOLO v3?

WATCH [THIS TED TALK](#) BY JOSEPH REDMON FOR AN OVERVIEW

DARKNET

DARKNET IS AN OPEN SOURCE
**NEURAL NETWORK
FRAMEWORK** WRITTEN IN C
AND CUDA.

IT IS FAST, EASY TO INSTALL,
AND SUPPORTS **CPU AND GPU**
COMPUTATION.

YOLO v3

YOU ONLY LOOK ONCE (YOLO)
IS A STATE-OF-THE-ART,
REAL-TIME **OBJECT DETECTION
SYSTEM**.



IT TOOK ABOUT **1.5 HOURS** TO TRAIN THE MODEL AND OBTAIN A “WEIGHTS” FILE

```

 7 max      2 x 2 / 2    52 x 52 x 128  ->  26 x 26 x 128
 8 conv    256 3 x 3 / 1    26 x 26 x 128  ->  26 x 26 x 256  0.399 BFLOPs
 9 max      2 x 2 / 2    26 x 26 x 256  ->  13 x 13 x 256
10 conv    512 3 x 3 / 1    13 x 13 x 256  ->  13 x 13 x 512  0.399 BFLOPs
11 max      2 x 2 / 1    13 x 13 x 512  ->  13 x 13 x 512
12 conv   1024 3 x 3 / 1    13 x 13 x 512  ->  13 x 13 x1024  1.595 BFLOPs
13 conv    256 1 x 1 / 1    13 x 13 x1024  ->  13 x 13 x 256  0.089 BFLOPs
14 conv    512 3 x 3 / 1    13 x 13 x 256  ->  13 x 13 x 512  0.399 BFLOPs
15 conv     18 1 x 1 / 1    13 x 13 x 512  ->  13 x 13 x 18   0.003 BFLOPs
16 yolo
17 route   13
18 conv    128 1 x 1 / 1    13 x 13 x 256  ->  13 x 13 x 128  0.011 BFLOPs
19 upsample      2x    13 x 13 x 128  ->  26 x 26 x 128
20 route   19 8
21 conv    256 3 x 3 / 1    26 x 26 x 384  ->  26 x 26 x 256  1.196 BFLOPs
22 conv     18 1 x 1 / 1    26 x 26 x 256  ->  26 x 26 x 18   0.006 BFLOPs
23 yolo

```

Loading weights from nalu_bu/yolov3-tiny-obj.backup...Done!

Learning Rate: 0.001, Momentum: 0.9, Decay: 0.0005

Resizing

576

Loaded: 0.000061 seconds



— STEP 2 —

WEB SERVER + DETECTION

Every 10 seconds the client takes a picture & POSTs to localhost:3000


**output = subprocess.getoutput("./darknet detect cfg/naluv1_tiny_net.cfg
naluv1_tiny.weights data/" + filename + " -thresh 0.4")**

count = output.count('nalu: ')

os.remove(app.config['UPLOAD_FOLDER'] + "/" + filename)

return jsonify({"count": count})

Python
Code



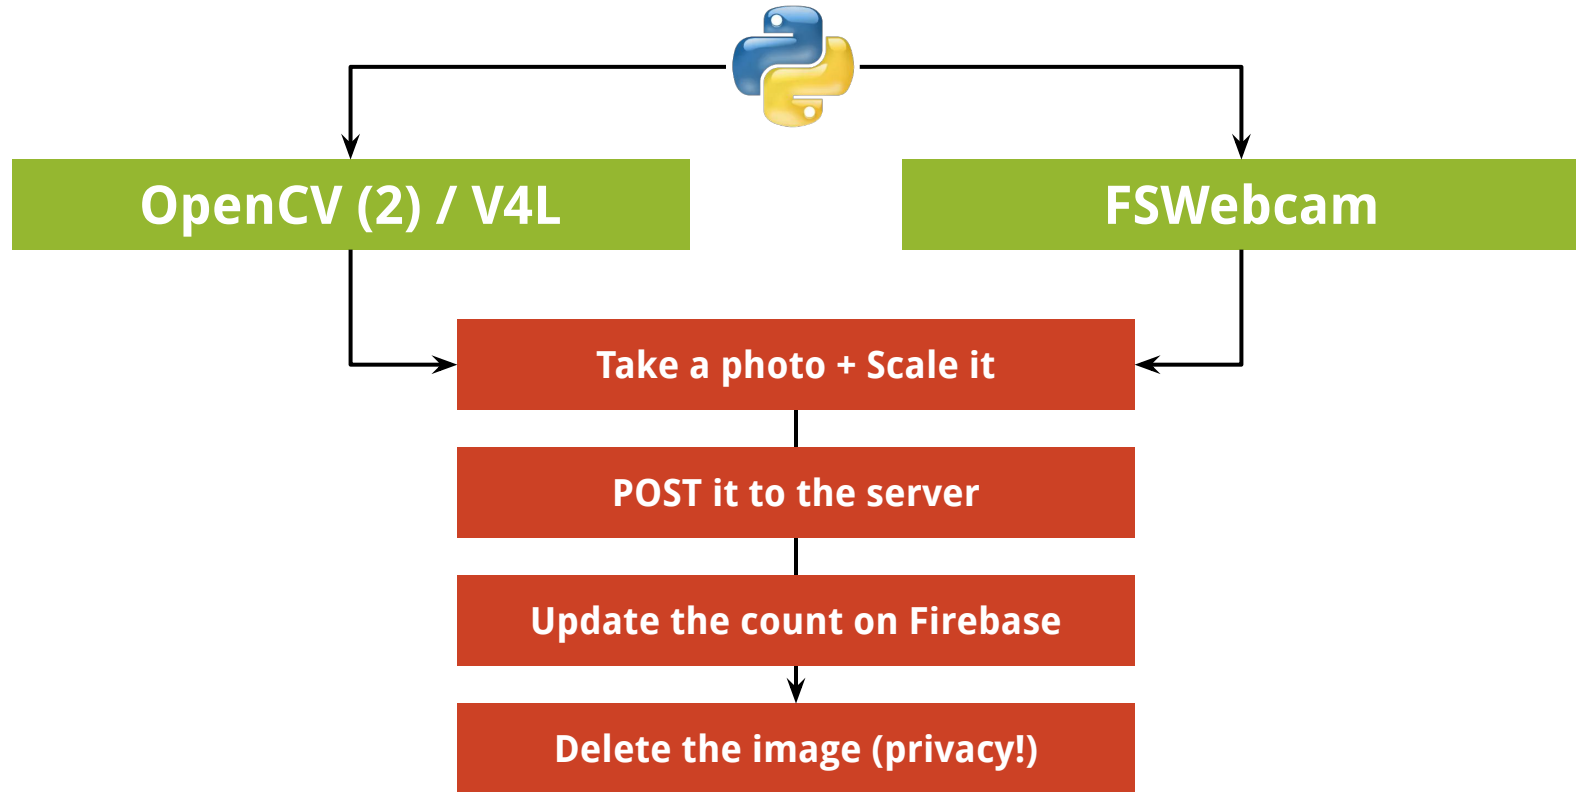
Flask

WE USED A SIMPLE FLASK SERVER (SYNC/BLOCKING)
TO AVOID OVERSATURATING DARKNET

— STEP 3 —

THE CLIENT

IT'S A SIMPLE PYTHON SCRIPT WHICH CAN RUN ON A MAC WITH OPENCV, OR ON LINUX






NOT PERFECT, BUT GOOD ENOUGH...

— STEP 4 —

THE FRONTEND

Simple frontend (imitating a fridge) written in one HTML file with a sprinkling of jQuery (for nostalgia)





**LOOKS MORE LIKE A
FRIDGE ON MOBILE!**

SOURCE CODE

<https://github.com/lemiffe/nalu-net/>

(The source code is crude as this started out as a simple hack to demonstrate it was possible)

RESULTS

We figured out after installing it that the current model does not work well with **Nalus that are over 4-5 meters away**, or if they are behind the glass doors of a refrigerator.

The model can be re-trained to recognise Nalus in more scenarios if needed. It requires more training data, ideally over 500 photos. The model currently has an **accuracy of 70-80%** depending on the resolution, lighting conditions, and the threshold.

THE END!

DRINK  RESPONSIBLY